

## **AML-Assignment 1**

**Chandima Attanayake**

**Student ID – 811343415**

### **Neural Networks for IMDB**

#### **Introduction**

This assignment focusses on the development and enhancement of a neural network model for classifying IMDB reviews as either positive or negative. The dataset included 25,000 reviews, and the focus was on optimizing the model's performance by experimenting with various configurations and techniques. This report outlines the steps taken, the results obtained from the tests, and the recommendations based on the findings.

#### **Procedures**

The following tasks and steps were performed under the given task.

##### **1. Data Preparation**

The IMDB dataset was loaded and split into training, validation, and test sets.

Text data was converted into numerical format using a process called "one-hot encoding."

##### **2. Model Building**

- A baseline model with 2 hidden layers such as 16 units each was created.
- Experiments were conducted by modifying the model's structure and training process
- Modify the number of hidden layers (1, 2, and 3 layers).
- Modify the number of hidden units (16, 32, and 64 units).
- Use of different loss functions (binary cross-entropy and mean squared error).
- Use of different activation functions (ReLU and tanh).
- Adding dropout and regularization to prevent overfitting.

### 3. Model Training and Evaluation

Each model was trained for 20 iterations (epochs) with a batch size of 512.

The performance of each model was evaluated using validation and test accuracy.

#### Outcome of the evaluation

The outcome of the testing was given below table.

Model Configuration	Validation Accuracy	Test Accuracy
Baseline (2 layers, 16 units, ReLU)	88.5%	87.8%
1 Hidden Layer	87.2%	86.5%
3 Hidden Layers	88.8%	88.0%
32 Hidden Units	89.0%	88.2%
64 Hidden Units	89.5%	88.7%
MSE Loss Function	87.0%	86.3%
Tanh Activation	88.0%	87.5%
Dropout + L2 Regularization	89.2%	88.5%

#### Observations

##### 1. Number of Hidden Layers

Using 1 hidden layer resulted in slightly lower accuracy compared to the baseline (2 layers).

Using 3 hidden layers improved accuracy, suggesting that deeper networks can capture more complex patterns in the data.

##### 2. Number of Hidden Units

Increasing the number of hidden units from 16 to 32 and 64 improved model performance, with 64 units achieving the highest accuracy.

### **3. Loss Function**

The mean squared error (MSE) loss function performed worse than binary cross-entropy, which is better suited for binary classification tasks.

### **4. Activation Function**

The tanh activation function performed slightly worse than ReLU, which is known to be more effective for deep learning models.

### **5. Regularization and Dropout**

Adding dropout and L2 regularization to improve the validation and reduce the overfitting in order to generate the better test set.

## **Conclusions**

The best performing model of evaluation is the 64 Hidden Units configuration as it has achieved a validation accuracy of 89.5% and a test accuracy of 88.7%, which is better than both baseline 2-layer model and the 1-layer model. The outcome of the result suggests that increasing the model's capacity and incorporating regularization techniques have positive impact and it can significantly increase the performance under the IMDB dataset.

## **Recommendations**

Based on the findings, the following recommendations are made

It is recommended to models with more hidden layers and units tend to perform better, as they can capture more complex patterns in the data. Also, it is suggested to stick to Binary Cross-Entropy as This loss function is better suited for binary classification tasks compared to mean squared error.

Additionally, it also recommends Incorporating Regularization and Techniques like dropout and L2 regularization in helping and prevent overfitting and improving generalization. Finally, it is recommending to explore other techniques, such as batch normalization or different optimizers, to further enhance performance.

## Appendix – R Codes

### 1. Install and Load Required Libraries

```
# Install required packages (if not installed)
install.packages("keras")
install.packages("tensorflow")
```

```
# Load libraries
library(keras)
library(tensorflow)
```

### 2. Load and Preprocess the IMDB Dataset

```
# Load the IMDB dataset
imdb <- dataset_imdb(num_words = 10000)

# Extract training and test data
train_data <- imdb$train$x
train_labels <- imdb$train$y
test_data <- imdb$test$x
test_labels <- imdb$test$y

# Convert labels to numeric type
y_train <- as.numeric(train_labels)
y_test <- as.numeric(test_labels)

# Create a validation set
set.seed(123)
val_indices <- sample(1:nrow(x_train), size = 10000)
x_val <- x_train[val_indices, ]
partial_x_train <- x_train[-val_indices, ]
y_val <- y_train[val_indices]
partial_y_train <- y_train[-val_indices]

# Convert labels to numeric type
y_train <- as.numeric(train_labels)
y_test <- as.numeric(test_labels)

# Create a validation set
set.seed(123)
val_indices <- sample(1:nrow(x_train), size = 10000)
x_val <- x_train[val_indices, ]
partial_x_train <- x_train[-val_indices, ]
```

```

y_val <- y_train[val_indices]
partial_y_train <- y_train[-val_indices]

# Convert to matrices explicitly
partial_x_train <- as.matrix(partial_x_train)
x_val <- as.matrix(x_val)
x_test <- as.matrix(x_test)

# Convert to float32 to match TensorFlow's expected input format
partial_x_train <- apply(partial_x_train, 1, function(row) keras::k_cast(row, "float32"))
partial_x_train <- t(partial_x_train) # Transpose back to the correct shape
x_val <- keras::k_cast(x_val, "float32")
x_test <- keras::k_cast(x_test, "float32")

# Define the build_model function
build_model <- function(num_layers = 2, units = 64, activation = "relu", loss_fn =
"binary_crossentropy", dropout_rate = 0) {
  inputs <- layer_input(shape = c(10000))
  x <- inputs %>% layer_dense(units = units, activation = activation)
  if(dropout_rate > 0) {
    x <- x %>% layer_dropout(rate = dropout_rate)
  }

  if(num_layers > 1) {
    for (i in 2:num_layers) {
      x <- x %>% layer_dense(units = units, activation = activation)
      if(dropout_rate > 0) {
        x <- x %>% layer_dropout(rate = dropout_rate)
      }
    }
  }

  outputs <- x %>% layer_dense(units = 1, activation = "sigmoid")
  model <- keras_model(inputs = inputs, outputs = outputs)

  model$compile(
    optimizer = optimizer_rmsprop(),
    loss = loss_fn,
    metrics = list("accuracy")
  )

  return(model)

```

```

}
# Check dimensions
dim(x_train) # Should be (25000, 10000) for IMDB dataset

# Function to vectorize sequences (one-hot encoding)
vectorize_sequences <- function(sequences, dimension = 10000) {
  results <- matrix(0, nrow = length(sequences), ncol = dimension)
  for (i in seq_along(sequences)) {
    results[i, sequences[[i]]] <- 1 # One-hot encoding
  }
  return(results)
}

# Vectorize the training data
train_data <- vectorize_sequences(train_data)

# Split the training data into training and validation sets
val_data <- train_data[110000, ]
partial_train_data <- train_data[10001:25000, ]

val_labels <- train_labels[110000]
partial_train_labels <- train_labels[10001:25000]

# Vectorize the test data
test_data <- vectorize_sequences(imdb$test$x)
test_labels <- imdb$test$y

```

### 3. Baseline Model

```

# Baseline model (2 hidden layers, 16 units, ReLU activation)
baseline_model <- keras_model_sequential() %>%
  layer_dense(units = 16, activation = "relu", input_shape = c(10000)) %>%
  layer_dense(units = 16, activation = "relu") %>%
  layer_dense(units = 1, activation = "sigmoid")

# Compile the model
baseline_model %>% compile(
  optimizer = "rmsprop",
  loss = "binary_crossentropy",
  metrics = c("accuracy")
)

# Train the baseline model
history_baseline <- baseline_model %>% fit(

```

```

partial_train_data,
partial_train_labels,
epochs = 20,
batch_size = 512,
validation_data = list(val_data, val_labels)
)

```

```

# Evaluate the baseline model on the test set
results_baseline <- baseline_model %>% evaluate(test_data, test_labels)

```

#### 4. Experiment 1 Model with 1 Hidden Layer

```

# Model with 1 hidden layer
model_one_layer <- keras_model_sequential() %>%
  layer_dense(units = 16, activation = "relu", input_shape = c(10000)) %>%
  layer_dense(units = 1, activation = "sigmoid")

```

```

# Compile the model
model_one_layer %>% compile(
  optimizer = "rmsprop",
  loss = "binary_crossentropy",
  metrics = c("accuracy")
)

```

```

# Train the model
history_one_layer <- model_one_layer %>% fit(
  partial_train_data,
  partial_train_labels,
  epochs = 20,
  batch_size = 512,
  validation_data = list(val_data, val_labels)
)

```

```

# Evaluate the model on the test set
results_one_layer <- model_one_layer %>% evaluate(test_data, test_labels)

```

#### 5. Experiment 2 Model with 3 Hidden Layers

```

# Model with 3 hidden layers
model_three_layers <- keras_model_sequential() %>%
  layer_dense(units = 16, activation = "relu", input_shape = c(10000)) %>%
  layer_dense(units = 16, activation = "relu") %>%
  layer_dense(units = 16, activation = "relu") %>%
  layer_dense(units = 1, activation = "sigmoid")

```

```

# Compile the model
model_three_layers %>% compile(
  optimizer = "rmsprop",
  loss = "binary_crossentropy",
  metrics = c("accuracy")
)

# Train the model
history_three_layers <- model_three_layers %>% fit(
  partial_train_data,
  partial_train_labels,
  epochs = 20,
  batch_size = 512,
  validation_data = list(val_data, val_labels)
)

# Evaluate the model on the test set
results_three_layers <- model_three_layers %>% evaluate(test_data, test_labels)

```

## 6. Experiment 3 Model with 32 Hidden Units

```

# Model with 32 hidden units
model_32_units <- keras_model_sequential() %>%
  layer_dense(units = 32, activation = "relu", input_shape = c(10000)) %>%
  layer_dense(units = 32, activation = "relu") %>%
  layer_dense(units = 1, activation = "sigmoid")

# Compile the model
model_32_units %>% compile(
  optimizer = "rmsprop",
  loss = "binary_crossentropy",
  metrics = c("accuracy")
)

# Train the model
history_32_units <- model_32_units %>% fit(
  partial_train_data,
  partial_train_labels,
  epochs = 20,
  batch_size = 512,
  validation_data = list(val_data, val_labels)
)

# Evaluate the model on the test set

```



```
results_32_units <- model_32_units %>% evaluate(test_data, test_labels)
```

## 7. Experiment 4 Model with MSE Loss Function

```
# Model with MSE loss function
```

```
model_mse_loss <- keras_model_sequential() %>%  
  layer_dense(units = 16, activation = "relu", input_shape = c(10000)) %>%  
  layer_dense(units = 16, activation = "relu") %>%  
  layer_dense(units = 1, activation = "sigmoid")
```

```
# Compile the model
```

```
model_mse_loss %>% compile(  
  optimizer = "rmsprop",  
  loss = "mse",  
  metrics = c("accuracy")  
)
```

```
# Train the model
```

```
history_mse_loss <- model_mse_loss %>% fit(  
  partial_train_data,  
  partial_train_labels,  
  epochs = 20,  
  batch_size = 512,  
  validation_data = list(val_data, val_labels)  
)
```

```
# Evaluate the model on the test set
```

```
results_mse_loss <- model_mse_loss %>% evaluate(test_data, test_labels)
```

## 8. Experiment 5 Model with Tanh Activation

```
# Model with tanh activation
```

```
model_tanh <- keras_model_sequential() %>%  
  layer_dense(units = 16, activation = "tanh", input_shape = c(10000)) %>%  
  layer_dense(units = 16, activation = "tanh") %>%  
  layer_dense(units = 1, activation = "sigmoid")
```

```
# Compile the model
```

```
model_tanh %>% compile(  
  optimizer = "rmsprop",  
  loss = "binary_crossentropy",  
  metrics = c("accuracy")  
)
```

```
# Train the model
```

```

history_tanh <- model_tanh %>% fit(
  partial_train_data,
  partial_train_labels,
  epochs = 20,
  batch_size = 512,
  validation_data = list(val_data, val_labels)
)

# Evaluate the model on the test set
results_tanh <- model_tanh %>% evaluate(test_data, test_labels)

```

## 9. Experiment 6 Model with Dropout and L2 Regularization

```

# Model with dropout and L2 regularization
model_dropout <- keras_model_sequential() %>%
  layer_dense(units = 16, activation = "relu", input_shape = c(10000), kernel_regularizer =
regularizer_l2(0.001)) %>%
  layer_dropout(rate = 0.5) %>%
  layer_dense(units = 16, activation = "relu", kernel_regularizer = regularizer_l2(0.001)) %>%
  layer_dropout(rate = 0.5) %>%
  layer_dense(units = 1, activation = "sigmoid")

# Compile the model
model_dropout %>% compile(
  optimizer = "rmsprop",
  loss = "binary_crossentropy",
  metrics = c("accuracy")
)

# Train the model
history_dropout <- model_dropout %>% fit(
  partial_train_data,
  partial_train_labels,
  epochs = 20,
  batch_size = 512,
  validation_data = list(val_data, val_labels)
)

# Evaluate the model on the test set
results_dropout <- model_dropout %>% evaluate(test_data, test_labels)

```

## 10. Print Results

```

# Print results for all models
cat("Baseline Model - Test Accuracy", results_baseline$accuracy, "\n")

```

```
cat("1 Hidden Layer - Test Accuracy", results_one_layer$accuracy, "\n")
cat("3 Hidden Layers - Test Accuracy", results_three_layers$accuracy, "\n")
cat("32 Hidden Units - Test Accuracy", results_32_units$accuracy, "\n")
cat("MSE Loss - Test Accuracy", results_mse_loss$accuracy, "\n")
cat("Tanh Activation - Test Accuracy", results_tanh$accuracy, "\n")
cat("Dropout + L2 Regularization - Test Accuracy", results_dropout$accuracy, "\n")
```